

Brand Kit OS × Leafpad

Toward Natively Compatible Products for Automated, On-Brand Blog Publishing

A proposal to make Brand Kit OS and Leafpad work together natively — pairing Leafpad's publishing engine with Brand Kit OS's editorial intelligence so users can generate consistently on-brand blog content end to end, with minimal manual effort.

Prepared for: The Leafpad Team

Prepared by: Brand Kit OS

Date: June 19, 2026

Re: Native MCP compatibility & shared field-mapping standard

1. Purpose & Framing

This is first and foremost a **partnership proposal**, not a feature-request list. Our goal is to build a plugin that lets users leverage the distinct strengths of **both** products — Leafpad's publishing and generation engine, and Brand Kit OS's brand intelligence — to produce automated, on-brand blog posts for their websites.

We have already built a working integration through the Model Context Protocol (MCP), and it works well today. This document explains *how* the two products complement each other, proposes a **shared field-mapping standard** so an AI agent can treat our data models as interoperable, and outlines a small set of **MCP exposures** that would take the combined experience from “works” to “native.”

THE ASK IN ONE LINE

Let's make Brand Kit OS and Leafpad **natively compatible**, so a user with both can go from brand strategy to published, on-brand blog post with a single automated pipeline.

2. Why the Two Products Fit Together

Leafpad's native understanding of a brand comes from its knowledge base, which retrieves over the organization's *existing published site content*. That makes Leafpad's generation **reactive** — it mirrors what a brand has already said. Powerful, but it has three structural limits:

- **Cold start:** new or thin sites have little content to imitate, so output drifts generic.
- **Drift compounding:** if existing content is inconsistent, generation learns and amplifies it.
- **No intended state:** retrieval captures what a brand has said — never what it should sound like, who it is for, or what it must never do.

Brand Kit OS supplies that missing layer: a structured source of truth for what a brand **intends to be** — voice, personality, governance, audience, products, and terminology. Together, the model becomes **prescriptive** instead of reactive.

What each product brings

Leafpad — the publishing body	Brand Kit OS — the editorial brain
Generation engine & scheduled topic generation	Prescriptive voice, tone & verbal style
Publishing surface, drafts, SEO fields, tags	Governance: forbidden terms, disclosures, formatting
Brand-styled AI image generation (palette auto-applied)	Audience & persona targeting
Knowledge base over existing site content (retrieval)	Product positioning, differentiators, competitive angle
Google Analytics integration & native analytics view	Preferred terminology & deliberate SEO lexicon

The complement, field by field

Leafpad gap (native)	Brand Kit OS fills it with
No explicit voice model — only imitation via retrieval	Voice archetypes, tone, verbal style, preferred terminology
No personality consistency across posts	Traits, values, principles, moods
No compliance / governance guardrails	Negative directory, disclosure policy, formatting constraints
No audience targeting	Audience personas & persona-aware framing
No accurate product positioning or CTAs	Products, differentiators, callouts

Leafpad gap (native)	Brand Kit OS fills it with
Freeform SEO keywords	Branded, deliberate keyword set & preferred terminology

3. How the Integration Works Today

Because Leafpad's generation runs server-side, Brand Kit OS adds value at the two points where the MCP exposes a seam. Both work now.

Leverage point 1 — Pre-generation brand brief

Leafpad's scheduled-post tool accepts an optional **prompt**. Brand Kit OS compiles a full brand brief — voice, governance rules, audience, preferred and forbidden terms — into that prompt, so even Leafpad's autonomous scheduled drafts come out on-brand instead of generic. This is the single highest-leverage integration: it upgrades the hands-off path, not just the manual one.

Leverage point 2 — Pre-publish QA on the draft

Every generated post lands as a **draft**. The agent reads the draft back, audits it against governance (forbidden terms, disclosures, formatting) and voice, and adjusts it before publish. Because this happens entirely in draft state, it is iteration — not remediation — with **no business-value loss**. The draft buffer is a safe staging area, and the user controls the publish gate.

NET EFFECT

Brand Kit OS becomes the standard a Leafpad blog is measured against — at generation time and again at publish time — giving users consistency at scale, across every post and every future scheduled post.

Strategic benefits this unlocks for Leafpad users

- **Cold-start solved:** brands with no site content still get on-brand output on day one.
- **Consistency over time:** one source of truth anchors every post against drift.
- **Multi-brand / agency leverage:** each brand kit is a distinct, swappable voice.
- **Compliance-grade content:** disclosures and forbidden-term enforcement make autonomous publishing safe for regulated brands.
- **Brand-aware imagery:** Leafpad applies palette and visual style automatically; Brand Kit OS supplies the on-brand subject and motif as the prompt.

4. Proposal: A Shared Field-Compatibility Protocol

The most valuable near-term step requires **no engineering on either side**: a shared reference that tells an AI agent *when a Brand Kit OS field and a Leafpad field are the same thing*, and what transform makes them interoperable. Today an agent knows what to send, but not the **compatibility semantics** behind each pairing. This protocol encodes that.

Each pairing carries one of five **match types**, so the agent knows how to treat it:

- **Exact** — same field, send as-is.
- **Transform** — same meaning, needs a deterministic conversion.
- **Partial** — related but not identical; map with a rule + caveat.
- **No-equivalent** — exists in one product only; drop or hold.
- **Auto-handled** — set server-side by Leafpad; the agent must not send it.

Brand Kit OS field	Leafpad field	Match type	Protocol for the agent
Drafted title	name	Exact	Send directly
Kebab-case slug	slug	Exact	Optional; Leafpad auto-derives if omitted
body_markdown	html_content	Transform	Convert MD → HTML; escape literal & in body
seo.title	seo_title	Exact (≤60)	Send directly; enforce 60-char cap
seo.description	seo_description	Exact (140–160)	Send raw text; do not HTML-escape
seo.keywords[] + preferred terms	seo_keywords	Transform	Join array → comma string
Content pillars / categories	tags	Partial	Map pillar → tag name; reuse existing; immutable after create
governance.disclosure	(in html_content)	Transform	Render disclosure into body HTML — no dedicated field
Persona / author	author	Auto-handled	Do not send; Leafpad sets byline. Persona informs voice
Visual style / palette	image styling	Auto-handled	Do not send; image tool applies brand style. Prompt = subject
excerpt	—	No-equivalent	Drop for Leafpad; retain for other channels
reading_time, canonical_url, visibility	—	No-equivalent	Drop on publish
Image alt / caption	— (requested)	No-equiv → pending	Embed in HTML today; see §5

WHY THIS MATTERS FOR THE PARTNERSHIP

When the two teams agree that a pair of fields *should* be treated as equivalent, we add a row and the agent immediately understands the compatibility — because the protocol column is an instruction, not just documentation. If Leafpad ships a schema change, one row updates and behavior follows. This is the foundation of “natively compatible.”

5. Where Native Support from Leafpad Would Help

These are the points where a small MCP exposure on Leafpad's side would remove friction and unlock a materially better joint experience. They are ordered by impact on the shared product.

#	Request	Why it helps the joint product
1	Schedule a finished post for future auto-publish (real published_at)	Today's scheduler generates a topic into a draft on a date; it cannot publish a pre-written post live on a schedule. A true publish date enables genuine editorial calendars.
2	Draft / status visibility — a status filter on the post-list tool and scheduled-queue read access	The agent can fetch a known draft by slug, but cannot enumerate drafts or see the scheduled queue. A status filter closes the QA and planning loop.
3	Image alt & caption as first-class write fields	We can generate brand-appropriate alt/caption text today, but it can only live in raw HTML. Write fields make accessibility and captions native and durable.
4	Mutable tags on update, and independently updatable SEO fields	Tags are immutable after creation (requires recreating the post), and SEO fields are all-or-nothing on update. Both force clumsy workarounds.
5	canonical_url field	Essential for syndicated / cross-posted content — a real SEO gap today.
6	Analytics read access via MCP (per-post views & GA metrics)	Leafpad already has the GA integration and native analytics; exposing it for read lets the brand strategy learn from what performs. Owned and best delivered by Leafpad.
7	Knowledge-base freshness / re-index control	The company-data tool is read-only Q&A with no way to trigger a re-index or check index freshness.

What we are deliberately not asking for

We do **not** want a delete or archive endpoint in the MCP. Manual-UI-only deletion is a feature, not a gap: the friction is intentional, and keeping destructive actions out of the automated path is a safeguard we want to preserve. We note this to show our requests are scoped and considered, not a wishlist.

6. Why This Is Good for Leafpad

- **Stickier output quality:** brand grounding raises the floor on every generated post, which reflects directly on Leafpad's perceived quality.
- **Expanded addressable use cases:** agencies and regulated brands become viable customers once governance and multi-brand voice are in the loop.
- **Differentiation:** “brand-governed, automated publishing” is a stronger story than generation alone — one neither product tells as well separately.
- **Low lift:** the compatibility protocol needs no engineering; the MCP requests are incremental field exposures, not new systems.

7. Proposed Next Steps

1. Align on the field-compatibility protocol (§4) as a shared, living reference both teams own.
2. Review the §5 MCP exposures and agree on a rough priority and sequence.
3. Define what “natively compatible” means as a milestone — the specific fields and flows that should work end to end without manual HTML or workarounds.
4. Pilot the combined pipeline on a shared test brand and review the output together.

IN SHORT

We are not trying to bolt onto Leafpad — we are proposing that our products meet in the middle, with a shared mapping standard and a handful of native fields, so a user with both can turn brand strategy into published, on-brand content automatically. We would value your team's buy-in to make that compatibility native.

Prepared by Brand Kit OS for discussion with the Leafpad team · June 19, 2026 · Field behaviors referenced here were calibrated against the live Leafpad MCP and are subject to re-verification on schema change.